

Mathematical Foundations for Power Analysis in Complex Systems

Gatlen Culp^{*1} Hamza Chaudhry² Herbie Bradley¹ Nandini Shiralkar¹

Abstract

Mathematical Framework for Power Analysis. Understanding societal risks—from institutional capture to AI-enabled power concentration—requires rigorous models of how power operates between agents with complex, evolving relationships. I develop a formal mathematical framework for power dynamics by extending Dahl’s influential 1957 theory, addressing its fundamental limitations through game-theoretic foundations.

Key Contributions. First, I demonstrate why absolute power models are insufficient, proving that relational power frameworks form a proper superset of absolute ones. Second, I show that Dahlian power cannot handle influence chains or compare power across different domains—critical limitations for analyzing complex systems. Third, I establish that utility functions are necessary for coherent power analysis, formalizing when power becomes behaviorally meaningful. This framework provides theoretical foundations for analyzing dangerous power dynamics in contexts where agents may lack clear utility functions or act irrationally over extended periods.

1. Introduction

Understanding power dynamics is critical for analyzing societal risks, from institutional capture to AI-enabled concentration of control. Yet despite decades of research across sociology, political science, and economics, we lack rigorous mathematical frameworks for modeling how power operates in complex multi-agent systems.

Consider these scenarios: A social media algorithm subtly shapes millions of political opinions through content curation. A powerful corporation influences regulatory policy

through a web of lobbying relationships. An AI system gradually assumes decision-making authority across critical infrastructure. Each involves intricate power relationships that existing theories struggle to model precisely.

The Challenge. Current approaches to power analysis suffer from fundamental limitations. Sociological theories offer rich conceptual frameworks but lack mathematical precision needed for quantitative analysis. Game-theoretic models provide formal rigor but assume rational agents with well-defined utility functions—assumptions that often fail in real-world power dynamics where agents act irrationally, preferences evolve, or utility functions remain unclear.

Our Approach. This paper develops a formal mathematical framework for power analysis by extending Dahl’s seminal 1957 theory of power. Dahl’s framework – “A has power over B to the extent that A can get B to do something B would not otherwise do”—provides an intuitive starting point but faces severe limitations when analyzing complex systems.

Contributions. We make three key theoretical contributions. First, we prove that models of relative power form a proper superset of absolute power models, establishing why relational frameworks are necessary. Second, we demonstrate that Dahlian Power cannot handle influence chains or enable cross-domain power comparisons – critical limitations for analyzing modern power structures. Third, we establish formal conditions under which utility functions become necessary for coherent power analysis.

Roadmap. §2 introduces our formal modeling approach and contextualizes it within existing power literature. §3 formalizes and extends Dahl’s framework, revealing its structural limitations. §4 proves the inadequacy of absolute power models. §5 demonstrates why utility functions are essential for meaningful power analysis. Together, these results provide theoretical foundations for analyzing dangerous power dynamics in complex systems.

^{*}Equal contribution ¹ERA:AI Research Fellowship ²Future of Life Institute. Correspondence to: Gatlen Culp <GatlenCulp@gmail.com>, Hamza Chaudhry <hamza@futureoflife.org>, Herbie Bradley <mail@herbiebradley.com>, Nandini Shiralkar <nandini@erafellowship.org>.

2. Towards a Formal Model of Power

The Literature Landscape. Power theory suffers from conceptual fragmentation. Existing approaches fall into two problematic categories: either loosely connected sociological definitions that resist formalization, or rigorous mathematical models confined to narrow domains (e.g., cooperative game theory’s power indices).¹ Each theoretical tradition critiques its predecessors while offering new definitions, creating a tower of competing frameworks rather than cumulative progress.²

Contemporary Applications. Formal power modeling has gained urgency across multiple domains. In reinforcement learning, researchers need to determine when agents pursue aligned objectives despite different behaviors and capabilities.³ AI governance researchers seek to formalize how labor-replacing AI disrupts social contracts and how economic power converts to political influence.⁴ While my framework may not immediately solve these challenges, it establishes necessary theoretical foundations.

Methodological Approach. Rather than proposing yet another universal “Theory of Power,” we develop rigorous mathematical foundations for one influential perspective – Dahl’s concept of power from 1957 – then systematically extend it. This approach allows us to either: (a) formally relate competing definitions to our framework, (b) demonstrate their mathematical incoherence, (c) identify

behaviorally irrelevant distinctions, or (d) establish approximation boundaries for complex realities.

As Dahl himself anticipated, we may never achieve “a single, consistent, coherent ‘Theory of Power’” but instead must develop “theories of limited scope” that prove useful within specific research contexts.⁵

The Value of Mathematical Formalization. Formal models provide conceptual precision beyond their mathematical machinery. When researchers invoke “prisoner’s dilemmas” or “principal-agent problems,” they rarely solve explicit equations – instead, they leverage rigorous conceptual frameworks that clarify otherwise murky situations. Similarly, our formal power framework aims to provide precise analytical tools for recognizing and categorizing power dynamics across diverse contexts. The mathematics enables conceptual clarity, not computational solutions.

Our Strategy. We begin with Dahl’s 1957 formalization – “*A* has power over *B* to the extent that *A* can get *B* to do something *B* would not otherwise do” – because it offers both intuitive appeal and mathematical tractability. We then systematically identify its limitations, demonstrate required extensions, and establish what mathematical properties any complete power theory must satisfy. This constructive approach reveals not just Dahl’s shortcomings, but fundamental constraints on power modeling itself.

¹This reflects a common trade-off in formal modeling: rigor versus generality. Cooperative game theory’s power indices (Shapley, Banzhaf, etc.) provide precise mathematical definitions but only apply within specific institutional contexts like weighted voting games  confirm. Conversely, sociological theories achieve broad applicability by sacrificing mathematical precision. The challenge is developing frameworks that maintain both formal rigor and sufficient generality to analyze diverse power structures across different domains.

²These sociological approaches capture different but related aspects of power, yet lack mathematical precision needed for quantitative analysis: Dahl’s behavioral definition focuses on observable influence over actions but cannot handle indirect influence chains or enable cross-domain power comparisons. French and Raven’s taxonomy (reward, coercive, legitimate, referent, expert power) describes power sources but provides no framework for measuring or comparing power magnitudes – it remains unclear whether these categories are irreducible or whether some (e.g., legitimate power) can be decomposed into others (reward/coercive mechanisms). Lukes’ three dimensions examine different levels of decision-making processes (overt conflict, covert agenda-setting, ideational preference-shaping) but offer descriptive taxonomies rather than predictive models. Actor-Network Theory treats power as emergent from network relationships but explicitly resists formal quantification, limiting its predictive capacity. While each tradition illuminates important facets—behavioral outcomes, psychological bases, structural processes, relational emergence—none provides the mathematical foundations necessary for rigorous analysis of complex, multi-level power dynamics.

³See Jason X, PhD student at Cambridge University

⁴See Liam Patell at GovAI

⁵Dahl’s Concept of Power 1957: “Thus we are not likely to produce certainly not for some considerable time to come anything like a single, consistent, coherent ‘Theory of Power.’ We are much more likely to produce a variety of theories of limited scope, each of which employs some definition of power that is useful in the context of the particular piece of research or theory but different in important respects from the definitions of other studies.”

3. Generalized Dahl Framework

3.1. Defining Dahlian Power

I start with Dahl's influential concept, which I term **Dahlian Power** to distinguish it from other power definitions we'll encounter (Figure 1).

Remark 3.1. (Dahl's Intuition) Consider two scenarios from the original essay⁶

1. A person on a street corner “commanding” drivers to use the right side (which they already do)
2. A police officer directing traffic to turn when it would normally go straight.

Only the second demonstrates genuine power – the ability to cause behavior that wouldn't otherwise occur.




Figure 1. Basic Dahlian Power relationship: Agent A uses means W to influence Agent B's response x .

Definition 3.2. (Dahlian Power w/ 2-means)

Agent A 's Dahlian Power over B with respect to the response x by means w is:

$$M\left(\frac{A}{B} : W, x\right) := \Pr(B, x \mid A, w) - \Pr(B, x \mid A, \bar{w})$$

Where:

- $\Pr(B, x \mid A, w)$ is the probability that B takes action x given A took action w .
- $\bar{w}$ is A not taking action w and W is the unresolved action

Remark 3.3. (Key Components) Dahl's model is made up of the following components:

⁶The full quote from the Concept of Power (Dahl 1957): “Suppose I stand on a street corner and say to myself, ‘I command all automobile drivers on this street to drive on the right side of the road’; suppose further that all the drivers actually do as I ‘command’ them to do; still, most people will regard me as mentally ill if I insist that I have enough power over automobile drivers to compel them to use the right side of the road. On the other hand, suppose a policeman is standing in the middle of an intersection at which most traffic ordinarily moves ahead; he orders all traffic to turn right or left; the traffic moves as he orders it to do. Then it accords with what I conceive to be the bedrock idea of power to say that the policeman acting in this particular role evidently has the power to make automobile drivers turn right or left rather than go ahead. My intuitive idea of power, then, is something like this: A has power over B to the extent that he can get B to do something that B would not otherwise do.”

1. **Agents.** At least two agents each of which have binary **actions** (e.g. take action w or not)
2. **Connection.** Some mechanism $\mathcal{J}$ linking agents (unspecified)
3. **Probability.** Distribution over agent-action pairs

What this doesn't require:

1. **Utility.** Doesn't say what either agent prefers.
2. **Causal Relationship.** Dahl explicitly avoids requiring causal traces – only observable behavioral differences matter. However, does note causality is required for real power, hence the means must precede the response temporally.
3. **Practicality.** Doesn't say whether the controller actually would or would not take action w in practice.




Figure 2. Environmental power: Agent actions influence environmental states or outcomes.

Remark 3.4. (Dahlian Power Over Environment) Dahl's framework extends naturally beyond human responders. Since the mathematical formalism only requires probabilistic responses $\Pr(x \mid A, w)$, we can analyze agent power over environmental outcomes. (Figure 2).

Letting $E :=$ environment:

$$M\left(\frac{A}{E} : W, x\right) := \max_{w^+ \in \mathcal{W}} \Pr(E = x \mid A, w^+) - \max_{w^- \in \mathcal{W}} \Pr(E = x \mid A, w^-)$$

Example: A company's carbon emissions policy W influences climate outcomes x (temperature rise).

This captures intuitive notions of “environmental power” while maintaining Dahl's rigorous probabilistic foundation.

3.2. Generalizing Beyond Binary Means

Dahl's binary framework (action taken or not) is limiting. Real agents choose from multiple options.




Figure 3. Multi-controller dynamics: Multiple agents $A_1, \dots, A_n$ simultaneously influence responder B .

Definition 3.1. (*Dahlian Power w/ n -means*) Let action space $\mathcal{W} := \{w_1, \dots, w_n \mid n \in \mathbb{N}^+\}$. Then:

$$M\left(\frac{A}{B} : W, x\right) := \max_{w^+ \in \mathcal{W}} \Pr(B, x \mid A, w^+) - \max_{w^- \in \mathcal{W}} \Pr(B, x \mid A, w^-)$$

This measures A 's maximum influence over B 's likelihood of taking action x .

Lemma 3.2. When $|\mathcal{W}| = 1$ (no choice), Dahlian Power $M = 0$, fitting our intuition that power requires agency.

... I think Dahl actually notes that you need similar means as well. Look back into this.

3.3. Multi-Agent Dahlian-Power

Real power dynamics rarely involve just one controller. Consider the example from Dahl's original paper where multiple legislators influence the outcome of a policy being passed or not (Figure 3):

When multiple controllers $\{A_1, \dots, A_n\}$ simultaneously influence the responder B , we need to extend our framework. Each agent A_i can take actions from their action space $\mathcal{W}_i$, creating a joint action profile $w = (w_1, \dots, w_n)$.

Dahl provides two comparative measures for ranking controllers with respect to a specific outcome x :




Figure 4. Marginal power analysis: Agent A_1 is focal (blue), while other agents $A_2, \dots, A_n$ are marginalized (gray).

Definition 3.1. (*Marginal Dahlian Power*)

Let $\mathcal{X} :=$ Set of possible responses, then:

The marginal power of A_i in optimizing for a specific response $x \in \mathcal{X}$ relative to a control condition w_i^0 (baseline action) is obtained by marginalizing over the expected actions of other controllers:⁷

$$\begin{aligned} M_+(A_i) &:= \max_{w_i \in \mathcal{W}_i} \mathbb{E}_{w_{-i} \sim \pi_{-i}} [\Pr(x \mid A_i, w_i, w_{-i})] \\ &\quad - \mathbb{E}_{w_{-i} \sim \pi_{-i}} [\Pr(x \mid A_i, w_i^0, w_{-i})] \\ M_-(A_i) &:= \max_{w_i \in \mathcal{W}_i} \mathbb{E}_{w_{-i} \sim \pi_{-i}} [\Pr(\bar{x} \mid A_i, w_i, w_{-i})] \\ &\quad - \mathbb{E}_{w_{-i} \sim \pi_{-i}} [\Pr(\bar{x} \mid A_i, w_i^0, w_{-i})] \\ M^*(A_i) &:= M_+(A_i) + M_-(A_i) \end{aligned}$$

where π_{-i} represents the distribution over other agents' action profiles $w_{-i} = (w_1, \dots, w_{i-1}, w_{i+1}, \dots, w_n)$, and w_i^0 is agent A_i 's baseline or "do nothing" action that serves as the control condition.⁸

This measures agent A_i 's maximum influence over outcome x when accounting for uncertainty in other agents' behavior.⁹

⁷This is not the exact definition that Dahl provides, but a more general one.

⁸If you don't care about M_- or M_+ , it may make more sense to omit a control action and just focus on M^* . This may as well be the case if only two actions are available and a control isn't well defined.

⁹Dahl mentions the **Problem of the Chameleon** in his original paper to describe the effect that A_i exerts the most power when they take the actions maximizing/minimizing $\Pr(x \mid A_i, w_i)$. This was called the problem of the chameleon because, in Dahl's original problem, he was looking at senators and passing legislation – senators could maximize their measured power by mirroring how everyone else voted. There is an empirical problem he encountered with collecting




Figure 5. Pairwise-conflict power analysis: Agent A_1 is focal (blue), A_2 is adversarial (red), while other agents $A_3, \dots, A_n$ are marginalized (gray).

Definition 3.2. (Pairwise-Conflict Dahlian Power)

For pairwise comparison between agents A_i and A_j regarding outcome x , where A_j optimizes against A_i while marginalizing over all other agents:

$$M_+''(A_i) := \max_{w_i \in \mathcal{W}_i} \min_{w_j \in \mathcal{W}_j} \mathbb{E}_{\mathbf{w}_{-i,-j} \sim \pi_{-i,-j}} \left[\Pr(x \mid A_i, w_i, A_j, w_j, \mathbf{w}_{-i,-j}) \right]$$

$$M_-''(A_i) := \max_{w_i \in \mathcal{W}_i} \min_{w_j \in \mathcal{W}_j} \mathbb{E}_{\mathbf{w}_{-i,-j} \sim \pi_{-i,-j}} \left[\Pr(\bar{x} \mid A_i, w_i, A_j, w_j, \mathbf{w}_{-i,-j}) \right]$$

$$M''(A_i) := M_+''(A_i) + M_-''(A_i)$$

where $\pi_{-i,-j}$ represents the distribution over all other agents' action profiles, and agent A_i maximizes their influence while agent A_j minimizes it.

Agent A_i has greater adversarial power than A_j if:

$$M''(A_i) > M''(A_j)$$

This measures agent A_i 's maximum sway over outcome x when facing optimal opposition from agent A_j .¹⁰

data in this way, but the theoretical issue with comparative Dahlian Power and its indifference towards preferences will be adjusted later.

¹⁰Dahl used this pairwise power to form a rough ranking of power among legislators from limited voting data.

It's reasonable to assume coalitions could form between our controllers, expanding their power beyond just a single agent. We can replace the original controllers with the new coalition as a singular controller, redefining the means of the coalition to be a combination of the original means.




Figure 6. Coalition power analysis: Agents A_1, A_2 collaborate (matching blue), while $A_3, \dots, A_n$ remain independent (gray).

Definition 3.3. (Combined Dahlian Power)

For coalition of controllers (referring to their indices) $C \subseteq \{1, \dots, n\}$ with joint action space $\mathcal{W}_C = \times_{i \in C} \mathcal{W}_i$:

$$M_+(C) := \max_{\mathbf{w}_C \in \mathcal{W}_C} \mathbb{E}_{\mathbf{w}_{-C} \sim \pi_{-C}} [\Pr(x \mid C, \mathbf{w}_C, \mathbf{w}_{-C})]$$

$$M_-(C) := \max_{\mathbf{w}_C \in \mathcal{W}_C} \mathbb{E}_{\mathbf{w}_{-C} \sim \pi_{-C}} [\Pr(\bar{x} \mid C, \mathbf{w}_C, \mathbf{w}_{-C})]$$

$$M^*(C) := M_+(C) + M_-(C)$$

where $\mathbf{w}_C = (w_i \mid \forall i \in C)$ is the coalition's joint action and π_{-C} represents the distribution over non-coalition agents' actions.

This measures the collective power of coordinating agents when they can synchronize their actions optimally.¹¹

Theorem 3.4. (Combinatorial Complexity of Power Comparisons) For n controllers, each can be assigned to: (1) coalition C_1 , (2) coalition C_2 , or (3) marginalized. This yields $\frac{3^n - 1}{2}$ distinct power comparisons.

Proof: Each controller has 3 assignments, giving 3^n total configurations. We subtract 1 degenerate case (all marginalized). The grand coalition (all in C_1 or all in C_2) is valid.

¹¹This formulation has many similarities with Cooperative Game Theory.

Since C_1 vs C_2 comparisons are symmetric, we divide by 2. Note: $(3^n - 1)$ is always even since 3^n is odd for all n , ensuring integer results.

Example: For $n = 3$ controllers: $\frac{3^3 - 1}{2} = \frac{26}{2} = 13$ comparisons, including grand coalition A_1, A_2, A_3 .

3.4. Environmental Power




Figure 7. Reverse environmental influence: Environment affects agent behavior through probabilistic states.

Remark 3.1. (The Environment as Controller) Can environments exert “power”? While environments don’t take deliberate actions, environmental states do influence agent behavior (Figure 7):

Example. Weather “determines” whether others will attend your picnic – rain reduces attendance probability.

While this stretches our definition of power depending on how you view the environment, it may still be sensible to define the environment as a virtual agent where possible environmental states become the action space.

Remark 3.2. (Modeling Agency as a Methodological Choice) Dahlian Power doesn’t require desires or utility functions – only the ability to select from multiple actions. This creates a philosophical tension:

- Deterministic view: If environments (or agents) have no genuine choice, then $|\mathcal{W}| = 1$, yielding zero Dahlian Power
- Probabilistic view: We assign probabilities to environmental states due to our uncertainty, treating nature as if it “chooses” stochastically

Practical Resolution. Whether discussing weather patterns or legislative decisions, we model both agents and environments probabilistically because:

- Complete determinism is computationally intractable
- Uncertainty is inherent to our observations
- Probabilistic models yield useful predictions

Whether you model environment/people as probabilistic behavior or an agent is a methodological choice, not an ontological claim about free will or determinism. We use whichever representation – agent or environment – best serves our analytical goals.

3.5. Means-Response Chains and the Limits of Dahlian Power




Figure 8. Means-response chain: Sequential influence propagation through intermediate agents.

Remark 3.1. (Means-Response Chains) Consider a means-response chain where A influences B , who influences C , and so on (Figure 8):

Remark 3.2. (The Attribution Problem) Dahlian Power fundamentally cannot handle indirect influence chains. To see why, consider:

- $A^{(0)}$ has Dahlian Power over $A^{(1)}$ ’s action $W^{(1)}$
- $A^{(1)}$ has Dahlian Power over $A^{(2)}$ ’s action $W^{(2)}$
- But what is $A^{(0)}$ ’s power over $A^{(2)}$ ’s action? This could theoretically be determined, yet bears no real relationship to the metrics above. For example: If $A^{(2)}$ has complete ($M = 1$) power over x , then what does any non-zero measure of power over x from $A^{(0)}$ or $A^{(1)}$ mean?¹²




Figure 9. Black-box abstraction: Intermediate chain collapsed into environmental probability function.

Proposition 3.3. (The Black-Box Abstraction Necessity) To analyze $A^{(0)}$ ’s influence on the final outcome x , Dahl’s framework forces us to collapse the entire intermediate chain into an environmental black box abstraction – ie: assuming that $A^{(0)}$ is a kind of “initial mover”, the only one capable of independent action while the subsequent agents take actions dependent on $A^{(0)}$ (Figure 9).

Remark 3.4. (Information Loss) This abstraction loses all structural information about:

- How influence propagates through the network
- Which intermediate agents have veto power
- Where bottlenecks or amplifications occur
- How power dissipates or concentrates along the chain

Remark 3.5. (Lukes’ Critique Formalized) While Lukes’ full critique goes beyond just this issue, this limitation forms one of Lukes’ criticisms of Dahlian Power. In studying legislative power, Dahl focused on observable voting behavior, missing upstream influences (Figure 10).

¹²This issue has some semblance to the issue of committing a crime under duress in law.




Figure 10. Luke's critique: Upstream power influences that Dahlian analysis misses in legislative contexts.

Remark 3.6. (Power Types Missing from Dahl) Luke identified two types of power Dahl's framework cannot capture:

1. **Bargaining Power:** Shaping which policies reach the legislative floor
2. **Ideational Power:** Determining which ideas are even consciously considered

These operate through influence chains that Dahlian Power must either ignore or compress into an opaque probability function.




Figure 11. Flattened chain representation: Sequential chain modeled as simultaneous multiple controllers.

Proposition 3.7. (Multiple Controllers as Partial Solution) We could remodel the chain as multiple simultaneous controllers (Figure 11).

Remark 3.8. (Limitations of Flattening) This flattening assumes all agents directly influence the outcome, ignoring:

- Hierarchical relationships (who influences whom)
- Sequential dependencies (temporal ordering matters)
- Conditional influences (power activated only under certain conditions)

The incapability to model these complex relationships is a serious limitation of the model. Luke's complaint of Dahl's empirical approach is, in part, a complaint that Dahl's power **lacks so much predictive power** by this inability as to be nearly useless. Any serious model of power must move beyond the limitations of this model.

Remark 3.9. (The Technical vs. Structural Distinction) This limitation isn't fundamentally about computational complexity. Computer scientists have developed sophisticated inference algorithms for modeling intricate probabilistic relationships – Bayesian networks, Markov networks, loopy




Figure 12. Complex probabilistic structure: Bayesian network showing latent variables and dependencies hidden within Dahl's environmental abstraction.

belief propagation, and variational inference can handle arbitrary dependency structures within the environmental “black box abstraction.” See Figure 12 for an illustration of how a bayes net, starting from each agent's influences on latent variables (which can be thought of as “chains-of-effect”), could model how a final response changes due to some action by a controller.

Remark 3.10. (The Dual Role Problem) The deeper issue is **structural**: Dahlian Power's foundational assumption that root agents make decisions independently creates an irreconcilable tension with influence chains. In means-response chains, intermediate agents are simultaneously:

- **Responders** to upstream influence (requiring dependent decision-making)
- **Controllers** in their own right (requiring independent decision-making)

This dual role cannot be coherently modeled within Dahl's framework without either:

1. Abandoning the independence assumption (breaking the foundational logic)
2. Collapsing the chain structure (losing predictive information)

3.6. Scope Selection and Limitations




Figure 13. Scope selection limitation: Dahlian Power cannot meaningfully compare power across different response variables.

Remark 3.1. (Scope Limitation)

As discussed in the original essay, Dahlian Power is limited by scope limitation – ie: the response that you’re trying to measure (Figure 13). There is not a very sensible way of comparing the power between some response x and some response y .

Example. From Dahl’s original essay:

“With an average probability approaching one, I can induce each of 10 students to come to class for an examination on a Friday afternoon when they would otherwise prefer to make off for New York or Northampton. With its existing resources and techniques, the New Haven Police Department can prevent about half the students who park along the streets near my office from staying beyond the legal time limit. Which of us has the more power?”

— Dahl’s Concept of Power 1957

Remark 3.2. (The Importance Dimension) For a general theory of power, it is extremely limiting to lack the ability to compare power in two separate contexts. While the quote from above was meant to point out how comparing power is difficult and perhaps arbitrary, there are situations for which there is a clear distinction between who has more power in an intuitive sense.

Imagine legislators have different Dahlian Power when voting on two distinct topics: (A) Whether a social welfare program will be approved and (B) Whether town hall’s bike shed will be painted green or blue.

Intuitively, if legislator Alice has considerably more Dahlian Power over (A) than legislator Bob, no matter what Bob’s Dahlian Power over (B) is, we would still say that Alice has more power. *Our intuitive sense of power relies not only on the ability to influence outcomes, but also how much we care about those outcomes.* This will be revisited later, but for now will be left as a limitation of Dahlian Power.

3.7. General Modeling Difficulties

Practical Modeling Challenges. Dahlian Power has some modeling difficulties that exist across many domains but are nonetheless worth mentioning:

Remark 3.1. (Action Space Selection Problem) Dahlian Power doesn’t explain why some actions are included or excluded from the action space, offsetting difficulties to the modeler. E.g. If a legislator can burn down the legislative building instead of voting, guaranteeing the policy doesn’t get passed, does the legislator have lots of power? In one sense (the Dahlian Power sense), they do, but in another sense they don’t have more power than they did before be-

cause the probability they make this decision is essentially 0 – it’s up to the modeler to recognize this possibility and leave it out.

Additionally, it’s important to determine which actions are available. Partly from Lukes’ response, covert bargaining happens before reaching the voting floor. These actions over time and ability to bargain need to be added to the action/strategy space (e.g. By defining tuples of actions indicating actions over time, (bargain for, vote for)).

Remark 3.2. (Controller Selection Problem) Partially from Lukes’ response, it’s important to consider every source of influence that influences the final variable. i.e. It’s not just the legislators that determine whether a policy passes, they’re also influenced by society, their family, and others to different degrees.

3.8. Conclusion On Dahlian Power

Summary. Dahlian Power provides a rigorous probabilistic foundation for analyzing influence relationships. At its core, Dahl’s framework measures power as the difference in probability: $M(\frac{A}{B} : W, x) := \Pr(B, x \mid A, w) - \Pr(B, x \mid A, \bar{w})$. This captures the intuitive notion that A has power over B when A can make B do something B wouldn’t otherwise do.

Requirements. This newly generalized framework requires: (1) At least one controller with defined action spaces, (2) probabilistic relationships between means and responses, and (3) observable behavioral differences. Notably, Dahlian Power does **not** require utility functions, causal mechanisms, or rational decision-making – only that controllers (or environment) can choose between alternatives ($|\mathcal{W}_i| > 1$).

Utility. Dahlian Power proves most valuable in contexts with: simultaneous decision-making by multiple controllers, well-defined response variables, and direct influence relationships. The framework works well with empirical data as Dahl originally used it – ranking agents by their influence over specific outcomes through marginal, pairwise-conflict, and coalition power measures.

Fundamental Limitations. Despite its elegance, Dahlian Power faces four critical constraints: (1) **Means-response chain incompatibility** – indirect influence through intermediate agents cannot be coherently modeled without collapsing chains into environmental black boxes, losing structural information; (2) **Scope limitation** – no mechanism exists for comparing power across different response variables, preventing holistic power assessment; (3) **Action space arbitrariness** – the framework provides no guidance for determining which actions to include or exclude; (4)

Controller selection ambiguity – no systematic method for identifying all relevant sources of influence.

These limitations point toward the necessity of utility-based frameworks that can handle preference-weighted cross-domain comparisons and structured influence networks – topics we explore in subsequent sections.

4. Why absolute power Fails to Model power Dynamics

Should this only be possible once utility is introduced?

From cooperative game theory, we have weighted graphs, rule-based representations, and weighted-voting games.

4.1. The Modeling Asymmetry Thesis

Proposition 4.1. *Models of relative power ($\mathcal{R}$) form a proper superset of models of absolute power ($\mathcal{A}$), establishing that $\mathcal{R}$ is more expressive than $\mathcal{A}$ in modeling capacity.*

Dahl says that there must be a “connection” for there to be a power dynamic.

4.1.1. Formal Statement

Definition 4.2. (Power Model Spaces) Let:

- $\mathcal{A} = \{M :$
 $M \text{ models power as properties of agents}\}$
- $\mathcal{R} = \{M :$
 $M \text{ models power as relational}$
 $\text{properties between agents}\}$

Then: $\mathcal{A} \subset \mathcal{R}$ (proper subset relation)

4.1.2. Supporting Arguments

Theorem 4.3. (Reduction Principle) *Any absolute power model can be reformulated as a relative power model by introducing an environmental baseline.*

Formal Construction. Given absolute power $P_{\text{abs}} : \text{Agent} \rightarrow \mathbb{R}^+$, we can define:

$$P_{\text{rel}(a,e)} = P_{\text{abs}(a)} - P_{\text{abs}(e)}$$

where $P_{\text{abs}(e)}$ serves as a zero-point reference for environment e .

Replace with an actual proof from graph theory

Physical Analogy. This construction mirrors potential energy in physics, where absolute potential V is meaningless without a reference frame:

$$V_{\text{rel}} = V_{\text{abs}} - V_{\text{ground}}$$

Incompleteness of Absolute Power Models ($\mathcal{A} \subsetneq \mathcal{R}$).

Theorem 4.4. (Limitation Theorem) *Absolute power models cannot capture essential relational dynamics without introducing factors outside their framework.*

Demonstrate some examples like in my original, such as distance as one factor and just other seemingly complex

set of environmental interaction factors that seem almost impossible to resolve absolute powers between entities to relative powers in a sensible way. In physics potential can be mapped to potential energy using distance, unlikely we have anything nearly as similar.

Core Problem. The power relationship between agents A and B cannot be logically derived from their absolute powers alone. Power lacks transitivity – if A dominates B and B dominates C , this does not guarantee A dominates C . This non-transitivity demonstrates that power cannot be reduced to a single absolute measure.

Example. Consider a circular dominance structure (like rock-paper-scissors):

- Rock defeats Scissors
- Scissors defeats Paper
- Paper defeats Rock

No absolute power assignment P_{abs} can capture these relationships.

Conclusion. Since capturing these dynamics requires factors $\notin \mathcal{A}$, we have $\mathcal{A} \neq \mathcal{R}$.

improve

4.1.3. Proof of Proper Subset Relation

From §1.2.1: $\mathcal{A} \subseteq \mathcal{R}$

From §1.2.2: $\exists M \in \mathcal{R}$ such that $M \notin \mathcal{A}$

Therefore: $\mathcal{A} \subset \mathcal{R}$ (proper subset) $\square$

4.1.4. Power Over Environment as Logical Encapsulation of What is Meant By “Absolute Power”

fin

4.1.5. Conclusion

Absolute power alone is insufficient for modeling power. For the remainder of this work, we consider only relational power – power over the environment or power over other agents.

Note: This does NOT tell us that relative power is sufficient, only that it may be.

5. Why Utility is Necessary for Power Analysis

5.1. The Inseparability Thesis

Proposition 5.1. *Power is meaningless without considering agent utilities – these elements are fundamentally inseparable in any coherent power framework.*

NOTE: Will likely replace this utility statement.

5.1.1. Utility Alignment

Core Insight. Power is an incoherent concept under complete alignment of utilities.

Definition 5.2. (Utility Alignment) Let:

- $U_A, U_B : \mathcal{S} \rightarrow \mathbb{R}$ be utility functions for agents A and B
- $\mathcal{S}$ be the state space
- $P(A \rightarrow B)$ be the power of A over B

If $U_A(s) \propto U_B(s) \forall s \in \mathcal{S}$ (perfect alignment), then $P(A \rightarrow B)$ is behaviorally irrelevant.

Illustrative Example. Consider a human-loving agent who can travel to an island, harm anywhere from 1 to 1,000,000 innocent people, and then escape unnoticed. Does their power increase with the number they *could* harm? This capacity is irrelevant:

- The human-loving agent doesn't benefit from causing harm and would prefer not to
- The potential victims have no desire to be harmed

The mere existence of this action is meaningless without considering whether it would ever be exercised.

▢ Replace with better example.

▢ Include behavioral misalignment, ie: power over

5.1.2. Agent Boundary Definition

Theorem 5.3. (Aggregation Principle) *Alignment relationships determine when individual power can be meaningfully aggregated into a single coherent agent:*

$$\left[P_{\text{collective}} = \sum_i P_i \right] \Leftrightarrow [U_i \approx U_j \forall i, j]$$

Application. Humanity's collective power is discussable precisely because humans share sufficient utility alignment over key domains.

Connection to Sociology This is akin to *power-within* discussed in sociological literature.

5.1.3. Environmental Mediation

Extended Framework. When modeling multi-agent systems with environmental interactions:

- Environment E can be treated as an agent with its own dynamics
- Even without direct conflict between agents, they may compete for environmental resources
- Agents may inflict harm indirectly through environmental manipulation

Power becomes relevant when agents compete for control over shared environmental factors, even if their ultimate goals don't directly conflict.

▢ Also could say that their actions are some mixed strategy over a number of actions. Not sure if this is true but not a crucial point.

5.1.4. Agents in Shared State Space

Agents of course, are themselves components of state spaces and they may have preferences over the state of themselves and others. Often times they have preferences for their survival which must be modeled in the utility function as well.

5.1.5. Critique of Light-Cone Formulation

Orthogonality Scenario. Consider alien civilization Ω with utilities orthogonal to humanity H . Imagine Ω cares only about arranging neutrinos into smiley faces – an activity that has zero impact on human flourishing. Meanwhile, humans optimize matter and energy arrangements with no effect on neutrino patterns. (We can additionally imagine that neither agent meaningfully interferes with the other, perhaps humanity and these aliens never interact)

Even if Ω achieves perfect control over neutrino arrangements while humans have only partial control over matter/energy, it's meaningless to compare their "power over the environment." The concept of environmental control is incoherent without specifying *which aspects* of the environment matter.

Key Insight. Power must be defined over relevant state space – control over factors that actually matter to the agents involved:

$$P_{\text{relevant}(A)} = \int_{\mathcal{S}_A} |\nabla U_A(s)| ds$$

where $\mathcal{S}_A := \{s \in \mathcal{S} : \partial U_A / \partial s \neq 0\}$ represents the relevant state space.

Note that this relevant state space is inherently chosen in reference to some agent and thus is incoherent outside of some utility reference frame.

5.2. Synthesis


Theorem 5.1. (*Completeness Requirement*) *A complete theory of power requires (although may be insufficient):*


1. *Relational framework ($\mathcal{R} \supset \mathcal{A}$)*
2. *Behavioral specifications for all agents (possibly without explicit utilities)*
3. *Action space definitions $\mathcal{A}_i$*
4. *Relevance constraints on state space $\mathcal{S}_i \subset \mathcal{S}$*

Implication. Power analysis without game-theoretic foundations is fundamentally incomplete. $\square$

 Power of signaling. US vs Japan and communicating the existence of the nuclear bomb. US in some way doesn't have power over Japan if it doesn't benefit from either (A) Annihilating it or (B) Communicating its Annihilation.

References





In broad strokes, Dahl's theory of power is similar to the rich literature of markov nets in computer science literature. With a few differences:

- Taken from the perspective of actions that players take. (However these can be mapped onto priors in a markov net so long as there are no loops, which of course there are. Might not actually be an issue.)
- Control must be causal. Since this is near-impossible to prove in-full, it suffices to say that it's certainly not causal if the response happens before the action.

A few nice properties and intuitions can however can be carried over:

- Markov blankets (You can isolate control by)
- Can be mapped to time-series HMMs.

Notoriously very difficult. ☐ Maybe also show some existing CS problem to apply control in a new context. ☐ Also maybe see what other studies of power exist out there that looked at causal relationships in this way. Certainly not a new understanding of power, just one that wasn't looked at broadly.

Already we are at a point where we can define control in MDPs, one of the most fundamental structures in reinforcement learning and applicable to many scenarios. ☐ Can we apply this to some textbook MDP scenario? ☐ Not sure I should transition to MDPs without bringing in utility as a pre-requisite of power.